

Seam Carving

20602 Computer Science (Algorithms)

Final Project

Ivan Golovko DSBA'23

Motivation

- It's often desirable to resize images without preserving the aspect ratio
- Cropping leads to loss of information




Original image




Cropped image

Motivation

- It's often desirable to resize images without preserving the aspect ratio
- Rescaling distorts the objects in the picture



Original image




Rescaled image

Motivation

- It's often desirable to resize images without preserving the aspect ratio
- Solution – **seam carving** (a.k.a. “liquid rescaling”)



Original image



Seam carving result

Seam carving

- Content-aware image resizing algorithm, created in 2005 by scientist from Mitsubishi Electric Research Lab
- General idea – identify and remove the least important areas
- At each step remove continuous paths of pixels (seams) that avoid borders of significant object



Step 1: Energy function

- We want to assign to each pixel certain measure of importance
- Possible solution – gradient magnitude (used widely in image processing)
- Basically, norm of 2-dimensional gradient of the grayscale image
- However, other energy functions might be used (Laplacian, local entropy, etc.)

Gradient magnitude

1. Convert to grayscale
2. Calculate gradient as the following convolution:

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * I$$


$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * I$$

3. Calculate magnitude (L_2 - or L_1 -norm):

$$G = |G_x| + |G_y|$$

Step 1: Energy Function

Example of computing G_x (for G_y the mechanism is the same):



Step 1: Energy Function



$$|\frac{\partial}{\partial x} \mathbf{I}| + |\frac{\partial}{\partial y} \mathbf{I}|$$



Step 2: Seam with lowest energy


Problem can be formulated as a shortest path problem in an edge-weighted DAG with no negative weights as follows:

1. Graph is constructed as shown in the picture
2. Each edge going into “pixel” vertex has weight equal to energy of this pixel
3. Edges going to t have weight 0

Then **Dijkstra’s algorithm** can be applied

For an image with N pixels $V \sim N$, $E \sim 3N$, so running time would be $O(E \log V) = O(N \log N)$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0



Step 2: Seam with lowest energy

However, in this particular case it's by far not the best idea, and length of all shortest paths can be found via **simple DP algorithm** with $O(N)$ running time:

DP algorithm for Seam carving (part 1)

Let $d[i, j]$ be length of shortest path to pixel (i, j) , and $e[i, j]$ – its energy. Then:

1. $d[0, j] = e[0, j]$
2. $d[i, j] = e[i, j] + \min\{d[i-1, j-1], d[i-1, j], d[i-1, j+1]\}$ for all $i > 0$

In the end, length of shortest seam is $\min_j\{d[n, j]\}$, i.e. smallest value in last row

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0



Shortest paths matrix $d[i, j]$

0	1	3	3

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

$$1 + \min\{0, 1\}$$


Shortest paths matrix $d[i, j]$

0	1	3	3
1			

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

$$3 + \min\{0, 1, 3\}$$



Shortest paths matrix $d[i, j]$

0	1	3	3
1	3		

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

$$3 + \min\{1, 3, 3\}$$


Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

$$1 + \min\{3, 3\}$$



Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

$$3 + \min\{1, 3, 3\}$$


Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4			

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0



Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4			

Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4

Step 2: Seam with lowest energy


Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4

Length and endpoint
of "shortest" seam



Step 2: Seam with lowest energy

Once we have the paths matrix and the endpoint of the shortest seam, we can trace it back to the top recursively:

DP algorithm for Seam carving (part 2)

Let $d[i, j]$ be length of shortest path to pixel (i, j) and $p[i]$ – column index of pixel in the shortest seam in row i

1. $p[n] = \mathit{argmin}_j \{d[n, j]\}$
2. $p[i-1] = \mathit{argmin}_{s \in \{-1, 0, 1\}} \{d[n, p[i] + s]\}$ for all $i < n$

Step 2: Seam with lowest energy


Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4

Length and endpoint
of "shortest" seam



Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4



Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4



Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4



Step 2: Seam with lowest energy

Energy matrix $e[i, j]$

0	1	3	3
1	3	3	1
3	4	1	0
4	4	0	0

Seam to be removed

Shortest paths matrix $d[i, j]$

0	1	3	3
1	3	4	4
4	5	4	4
8	8	4	4


Results

Algorithm proceeds iteratively removing seams one by one, until the desired degree of compression is achieved. The process in progress looks [like this](#):



SC for image enlargement

- Seam carving algorithm can be easily adapted for enlarging images with similar idea – **insert copies of least important seams**
- However, a caveat is that if done straightforwardly, one seam will be inserted repeatedly, causing strange results:



SC for image enlargement

- Solution – first proceed like if we wanted to delete same number of seams
- Insert one copy of each seam that was removed in the first step



Extensions

- Very flexible, can work with different energy functions
- Can preserve the marked areas (assigns large costs)
- Can do object removal (remove seams until all the marked pixels disappear)
- However, doesn't always produce good-looking results



Some examples

Original image



Rescaled (x0.6)



Some examples

Seams to be deleted



SC-compressed (x0.6)



Some examples

Rescaled (x1.3)



Some examples

SC-enlarged (x1.3)




Some examples

Original image



Rescaled (x0.6)



Some examples

Seams to be deleted



SC-compressed (x0.6)



Some examples

SC-compressed (x0.6)



Rescaled (x0.6)



References

- Original paper by Avidan & Shamir:
<https://perso.crans.org/frenoy/matlab2012/seamcarving.pdf>
- GitHub repository with my implementation:
https://github.com/igolovko3/seam_carving

Discussion: Compressing in two dimensions


Assume we want to remove both vertical and horizontal seams from an image. A priori, it is not clear what should the order of removal be (horizontal first, vertical first, alternately, etc.)

Suppose $\mathbf{C}[i, j]$ – minimal cost of removing i horizontal and j vertical seams, $\mathbf{Im}[i, j]$ is the “optimal” image of corresponding size. Then:

1. $\mathbf{C}[0, j]$ and $\mathbf{C}[i, 0]$ are known for any i and j
2. $\mathbf{C}[i, j] = \min\{\begin{array}{l} \text{“Cost to remove one horizontal seam from } \mathbf{Im}[i - 1, j]\text{”} + \mathbf{C}[i - 1, j], \\ \text{“Cost to remove one vertical seam from } \mathbf{Im}[i, j - 1]\text{”} + \mathbf{C}[i, j - 1] \end{array}\}$
for all $i, j > 0$

Discussion: Topological sort VS DP

- Topological sorting – linear ordering of vertices of a DAG, such that if there is an edge $v \rightarrow u$, then v comes before u in that ordering
- Can be done in $O(V + E)$ by a modification of a DFS




Discussion: Topological sort VS DP

Can be applied to solve shortest path problem:

1. Get topological sort of graph – $O(E + V)$
2. Proceed like in Dijkstra's, but instead of searching for vertex with smallest distance to s on each step, just follow the topological order – $O(E + V)$


Correct, since any vertex v is unreachable from any of the following vertices, so shortest path to it doesn't depend on what happens after



Discussion: Topological sort VS DP

In the seam carving problem, graph is structured so that going over it **row by row** yields a **topologically-ordered traversal**, since all edges go “down”

Thus, start with step 2 right away




Discussion: Topological sort VS DP

- Let $d[v]$ be the length of shortest path from s to v
- After relaxing edges from A , B & C , $d[D]$ will be implicitly be set to:


$$d[D] = \min\{d[A] + |AD|, d[B] + |BD|, d[C] + |CD|\}$$

- However, since $|AD| = |BD| = |CD| = \text{energy}(D)$, this is the exact same expression as in DP algorithm
- They are **absolutely equivalent!**



Discussion: Fancy fact about seams

Though each seam is continuous on the image from which it is deleted, on the original image they might cross each other and “tear”



Discussion: Fancy fact about seams

- However, it can be shown that any set of number of seams obtained by iterative removals can be replaced by the same number of non-crossing continuous seams on the original image, covering same set of pixels
- Idea – each time take “left envelope” of the set of pixels, proceed iteratively

